

In [1]:

```
#importing libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import random
import seaborn as sns
from keras.preprocessing.image import ImageDataGenerator
```

In [2]:

```
#importing the dataset
train=pd.read_csv('sign_mnist_train.csv')
test=pd.read_csv('sign_mnist_test.csv')
```

In [3]:

```
print(train.shape)
print(test.shape)
```

(27455, 785)
(7172, 785)

In [4]:

```
train.head()
```

Out[4]:

	label	pixel1	pixel2	pixel3	pixel4	pixel5	pixel6	pixel7	pixel8	pixel9	...	pixel775	pixel776	pixel777	pixel778	pixel779	pixel780
0	3	107	118	127	134	139	143	146	150	153	...	207	207	207	207	206	207
1	6	155	157	156	156	156	157	156	158	158	...	69	149	128	87	94	16
2	2	187	188	188	187	187	186	187	188	187	...	202	201	200	199	198	19
3	2	211	211	212	212	211	210	211	210	210	...	235	234	233	231	230	22
4	13	164	167	170	172	176	179	180	184	185	...	92	105	105	108	133	16


5 rows x 785 columns

In [5]:

```
plt.figure(figsize = (10,10)) # Label Count
sns.set_style("darkgrid")
sns.countplot(train['label'])
```

Out[5]:

<matplotlib.axes._subplots.AxesSubplot at 0x25e377c9ee0>



In [7]:

```
# Create training and testing arrays
train_set = np.array(train, dtype = 'float32')
test_set = np.array(test, dtype='float32')
```

In [8]:


```
#Specifying class labels
class_names = ['A', 'B', 'C', 'D', 'E', 'F', 'G', 'H', 'I', 'J', 'K', 'L', 'M', 'N', 'O', 'P', 'Q', 'R', 'S', 'T', 'U', 'V', 'W', 'X', 'Y' ]
```

In [9]:

```
#See a random image for class label verification
i = random.randint(1,27455)
plt.imshow(train_set[i,1:].reshape((28,28)))
plt.imshow(train_set[i,1:].reshape((28,28)))
label_index = train["label"][i]
plt.title(f"{class_names[label_index]}")
plt.axis('off')
```


Out[9]:

(-0.5, 27.5, 27.5, -0.5)



In [10]:

```
# Define the dimensions of the plot grid
W_grid = 5
L_grid = 5
fig, axes = plt.subplots(L_grid, W_grid, figsize = (10,10))
axes = axes.ravel() # flatten the 15 x 15 matrix into 225 array
n_train = len(train_set) # get the length of the train dataset
# Select a random number from 0 to n_train
for i in np.arange(0, W_grid * L_grid): # create evenly spaces variables
    # Select a random number
    index = np.random.randint(0, n_train)
    # read and display an image with the selected index
    axes[i].imshow( train_set[index,1:].reshape((28,28)) )
    label_index = int(train_set[index,0])
    axes[i].set_title(class_names[label_index], fontsize = 8)
    axes[i].axis('off')
plt.subplots_adjust(hspace=0.4)
```



In [11]:

```
# Prepare the training and testing dataset
```

```
X_train = train_set[:, 1:] / 255
```

```
y_train = train_set[:, 0]
```

```
X_test = test_set[:, 1:] / 255
```

```
y_test = test_set[:,0]
```

In [12]:

```
#Visualize train images
```

```
plt.figure(figsize=(10, 10))
```

```
for i in range(25):
```

```
    plt.subplot(5, 5, i + 1)
```

```
    plt.xticks([])
```

```
    plt.yticks([])
```

```
    plt.grid(False)
```

```
    plt.imshow(X_train[i].reshape((28,28)), cmap=plt.cm.binary)
```

```
    label_index = int(y_train[i])
```

```
    plt.title(class_names[label_index])
```

```
plt.show()
```



In [13]:

```
#Split the training and test sets
```

```
from sklearn.model_selection import train_test_split
```

```
X_train, X_validate, y_train, y_validate = train_test_split(X_train, y_train, test_size = 0.2, random_state = 12345)
```

```
print(X_train.shape)
```

```
print(y_train.shape)
```

```
(21964, 784)
```

```
(21964,)
```

In [14]:

```
# Unpack the training and test tuple
```

```
X_train = X_train.reshape(X_train.shape[0], *(28, 28, 1))
```

```
X_test = X_test.reshape(X_test.shape[0], *(28, 28, 1))
```

```
X_validate = X_validate.reshape(X_validate.shape[0], *(28, 28, 1))
```

```
print(X_train.shape)
```

```
print(y_train.shape)
```

```
print(X_validate.shape)
```

```
(21964, 28, 28, 1)
```

```
(21964,)
```

```
(5491, 28, 28, 1)
```

In [15]:

```
# With data augmentation to prevent overfitting
datagen = ImageDataGenerator(
    featurewise_center=False, # set input mean to 0 over the dataset
    samplewise_center=False, # set each sample mean to 0
    featurewise_std_normalization=False, # divide inputs by std of the dataset
    samplewise_std_normalization=False, # divide each input by its std
    zca_whitening=False, # apply ZCA whitening
    rotation_range=10, # randomly rotate images in the range (degrees, 0 to 180)
    zoom_range = 0.1, # Randomly zoom image
    width_shift_range=0.1, # randomly shift images horizontally (fraction of total width)
    height_shift_range=0.1, # randomly shift images vertically (fraction of total height)
    horizontal_flip=False, # randomly flip images
    vertical_flip=False) # randomly flip images
datagen.fit(X_train)
```

In [16]:

```
#Library for CNN Model
import keras
from keras.models import Sequential
from keras.layers import Conv2D, MaxPooling2D, Dense, Flatten, Dropout, BatchNormalization
from keras.optimizers import Adam
from keras.callbacks import ReduceLROnPlateau
```

In [17]:

```
#Defining the Convolutional Neural Network
cnn_model = Sequential()
cnn_model.add(Conv2D(32, (3, 3), input_shape = (28,28,1), activation='relu'))
cnn_model.add(BatchNormalization())
cnn_model.add(MaxPooling2D(pool_size = (2, 2)))
cnn_model.add(Conv2D(64, (3, 3), input_shape = (28,28,1), activation='relu'))
cnn_model.add(Dropout(0.2))
cnn_model.add(BatchNormalization())
cnn_model.add(MaxPooling2D(pool_size = (2, 2)))
cnn_model.add(Conv2D(128, (3, 3), input_shape = (28,28,1), activation='relu'))
cnn_model.add(BatchNormalization())
cnn_model.add(MaxPooling2D(pool_size = (2, 2)))
cnn_model.add(Flatten())
cnn_model.add(Dense(units = 512, activation = 'relu'))
cnn_model.add(Dropout(0.25))
cnn_model.add(Dense(units = 25, activation = 'softmax'))
cnn_model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 26, 26, 32)	320
batch_normalization (Batch Normalization)	(None, 26, 26, 32)	128
max_pooling2d (MaxPooling2D)	(None, 13, 13, 32)	0
conv2d_1 (Conv2D)	(None, 11, 11, 64)	18496
dropout (Dropout)	(None, 11, 11, 64)	0
batch_normalization_1 (Batch Normalization)	(None, 11, 11, 64)	256
max_pooling2d_1 (MaxPooling2D)	(None, 5, 5, 64)	0
conv2d_2 (Conv2D)	(None, 3, 3, 128)	73856
batch_normalization_2 (Batch Normalization)	(None, 3, 3, 128)	512
max_pooling2d_2 (MaxPooling2D)	(None, 1, 1, 128)	0
flatten (Flatten)	(None, 128)	0
dense (Dense)	(None, 512)	66048
dropout_1 (Dropout)	(None, 512)	0
dense_1 (Dense)	(None, 25)	12825
Total params: 172,441		
Trainable params: 171,993		
Non-trainable params: 448		

In [18]:

```
#Compiling
cnn_model.compile(loss = 'sparse_categorical_crossentropy', optimizer='adam' ,metrics = ['accuracy'])
```

In [19]:

```
learning_rate_reduction=ReduceLRonPlateau(monitor='val_accuracy',patience=2,verbose=1,factor=0.5,min_lr=0.00001)
```

In [20]:

```
#Training the CNN model
history = cnn_model.fit(datagen.flow(X_train, y_train, batch_size = 512), epochs = 20, verbose=1, validation_data
= (X_validate, y_validate),callbacks=[learning_rate_reduction])
```

```
Epoch 1/20
43/43 [=====] - 37s 850ms/step - loss: 1.6502 - accuracy: 0.4982 - val_loss
: 3.2344 - val_accuracy: 0.0459
Epoch 2/20
43/43 [=====] - 36s 836ms/step - loss: 0.5078 - accuracy: 0.8281 - val_loss
: 3.2623 - val_accuracy: 0.0459
Epoch 3/20
43/43 [=====] - 35s 817ms/step - loss: 0.2523 - accuracy: 0.9173 - val_loss
: 3.2628 - val_accuracy: 0.0563
Epoch 4/20
43/43 [=====] - 35s 815ms/step - loss: 0.1578 - accuracy: 0.9491 - val_loss
: 3.2407 - val_accuracy: 0.0829
Epoch 5/20
43/43 [=====] - 35s 817ms/step - loss: 0.1102 - accuracy: 0.9651 - val_loss
: 3.0212 - val_accuracy: 0.1776
Epoch 6/20
43/43 [=====] - 35s 816ms/step - loss: 0.0844 - accuracy: 0.9734 - val_loss
: 2.9480 - val_accuracy: 0.1945
Epoch 7/20
43/43 [=====] - 35s 809ms/step - loss: 0.0660 - accuracy: 0.9788 - val_loss
: 2.9165 - val_accuracy: 0.1601
Epoch 8/20
43/43 [=====] - 35s 810ms/step - loss: 0.0578 - accuracy: 0.9810 - val_loss
: 2.1802 - val_accuracy: 0.3779
Epoch 9/20
43/43 [=====] - 35s 810ms/step - loss: 0.0518 - accuracy: 0.9838 - val_loss
: 1.4570 - val_accuracy: 0.5305
Epoch 10/20
43/43 [=====] - 36s 828ms/step - loss: 0.0448 - accuracy: 0.9861 - val_loss
: 0.9138 - val_accuracy: 0.7130
Epoch 11/20
43/43 [=====] - 35s 823ms/step - loss: 0.0388 - accuracy: 0.9870 - val_loss
: 0.3053 - val_accuracy: 0.9051
Epoch 12/20
43/43 [=====] - 35s 824ms/step - loss: 0.0342 - accuracy: 0.9889 - val_loss
: 0.2977 - val_accuracy: 0.8986
Epoch 13/20
43/43 [=====] - 35s 822ms/step - loss: 0.0306 - accuracy: 0.9910 - val_loss
: 0.0701 - val_accuracy: 0.9783
Epoch 14/20
43/43 [=====] - 34s 790ms/step - loss: 0.0312 - accuracy: 0.9909 - val_loss
: 0.0479 - val_accuracy: 0.9869
Epoch 15/20
43/43 [=====] - 35s 810ms/step - loss: 0.0259 - accuracy: 0.9914 - val_loss
: 0.0188 - val_accuracy: 0.9931
Epoch 16/20
43/43 [=====] - 34s 781ms/step - loss: 0.0260 - accuracy: 0.9919 - val_loss
: 0.0184 - val_accuracy: 0.9953
Epoch 17/20
43/43 [=====] - 34s 795ms/step - loss: 0.0240 - accuracy: 0.9927 - val_loss
: 0.1633 - val_accuracy: 0.9443
Epoch 18/20
43/43 [=====] - ETA: 0s - loss: 0.0218 - accuracy: 0.9920
Epoch 00018: ReduceLRonPlateau reducing learning rate to 0.0005000000237487257.
43/43 [=====] - 34s 797ms/step - loss: 0.0218 - accuracy: 0.9920 - val_loss
: 0.0313 - val_accuracy: 0.9867
Epoch 19/20
43/43 [=====] - 34s 798ms/step - loss: 0.0147 - accuracy: 0.9953 - val_loss
: 0.0022 - val_accuracy: 0.9995
Epoch 20/20
43/43 [=====] - 34s 801ms/step - loss: 0.0128 - accuracy: 0.9960 - val_loss
: 0.0018 - val_accuracy: 0.9995
```

In [21]:


```
#Visualizing the training performance
plt.figure(figsize=(12, 8))

plt.subplot(2, 2, 1)
plt.plot(history.history['loss'], label='Loss')
plt.plot(history.history['val_loss'], label='val_Loss')
plt.legend()
plt.grid()
plt.title('Loss evolution')

plt.subplot(2, 2, 2)
plt.plot(history.history['accuracy'], label='accuracy')
plt.plot(history.history['val_accuracy'], label='val_accuracy')
plt.legend()
plt.grid()
plt.title('Accuracy evolution')
```

Out[21]:

Text(0.5, 1.0, 'Accuracy evolution')



In [22]:

```
#Predictions for the test data
predicted_classes = cnn_model.predict_classes(X_test)
```

WARNING:tensorflow:From <ipython-input-22-2f08a57fd4c0>:2: Sequential.predict_classes (from tensorflow.python.keras.engine.sequential) is deprecated and will be removed after 2021-01-01.

Instructions for updating:

Please use instead: * `np.argmax(model.predict(x), axis=-1)`, if your model does multi-class classification (e.g. if it uses a `'softmax'` last-layer activation). * `(model.predict(x) > 0.5).astype("int32")`, if your model does binary classification (e.g. if it uses a `'sigmoid'` last-layer activation).

In [23]:

```
L = 5
W = 5
fig, axes = plt.subplots(L, W, figsize = (12,12))
axes = axes.ravel()

for i in np.arange(0, L * W):
    axes[i].imshow(X_test[i].reshape(28,28))
    axes[i].set_title(f"Prediction Class = {predicted_classes[i]:0.1f}\n True Class = {y_test[i]:0.1f}")
    axes[i].axis('off')
plt.subplots_adjust(wspace=0.5)
```



In [24]:

```
from sklearn.metrics import confusion_matrix
from sklearn import metrics
cm = metrics.confusion_matrix(y_test, predicted_classes)
```


In [25]:

```
#Defining function for confusion matrix plot
def plot_confusion_matrix(y_true, y_pred, classes,
                           normalize=False,
                           title=None,
                           cmap=plt.cm.Blues):

    if not title:
        if normalize:
            title = 'Normalized confusion matrix'
        else:
            title = 'Confusion matrix, without normalization'

    # Computing confusion matrix
    cm = confusion_matrix(y_true, y_pred)
    if normalize:
        cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
        print("Normalized confusion matrix")
    else:
        print('Confusion matrix, without normalization')

# Visualizing
    fig, ax = plt.subplots(figsize=(10, 10))
    im = ax.imshow(cm, interpolation='nearest', cmap=cmap)
    ax.figure.colorbar(im, ax=ax)
    # We want to show all ticks...
    ax.set(xticks=np.arange(cm.shape[1]),
           yticks=np.arange(cm.shape[0]),
           xticklabels=classes, yticklabels=classes,
           title=title,
           ylabel='True label',
           xlabel='Predicted label')

    # Rotating the tick labels and setting their alignment.
    plt.setp(ax.get_xticklabels(), rotation=45, ha="right",
              rotation_mode="anchor")

    # Looping over data dimensions and create text annotations.
    fmt = '.2f' if normalize else 'd'
    thresh = cm.max() / 2.
    for i in range(cm.shape[0]):
        for j in range(cm.shape[1]):
            ax.text(j, i, format(cm[i, j], fmt),
                    ha="center", va="center",
                    color="white" if cm[i, j] > thresh else "black")

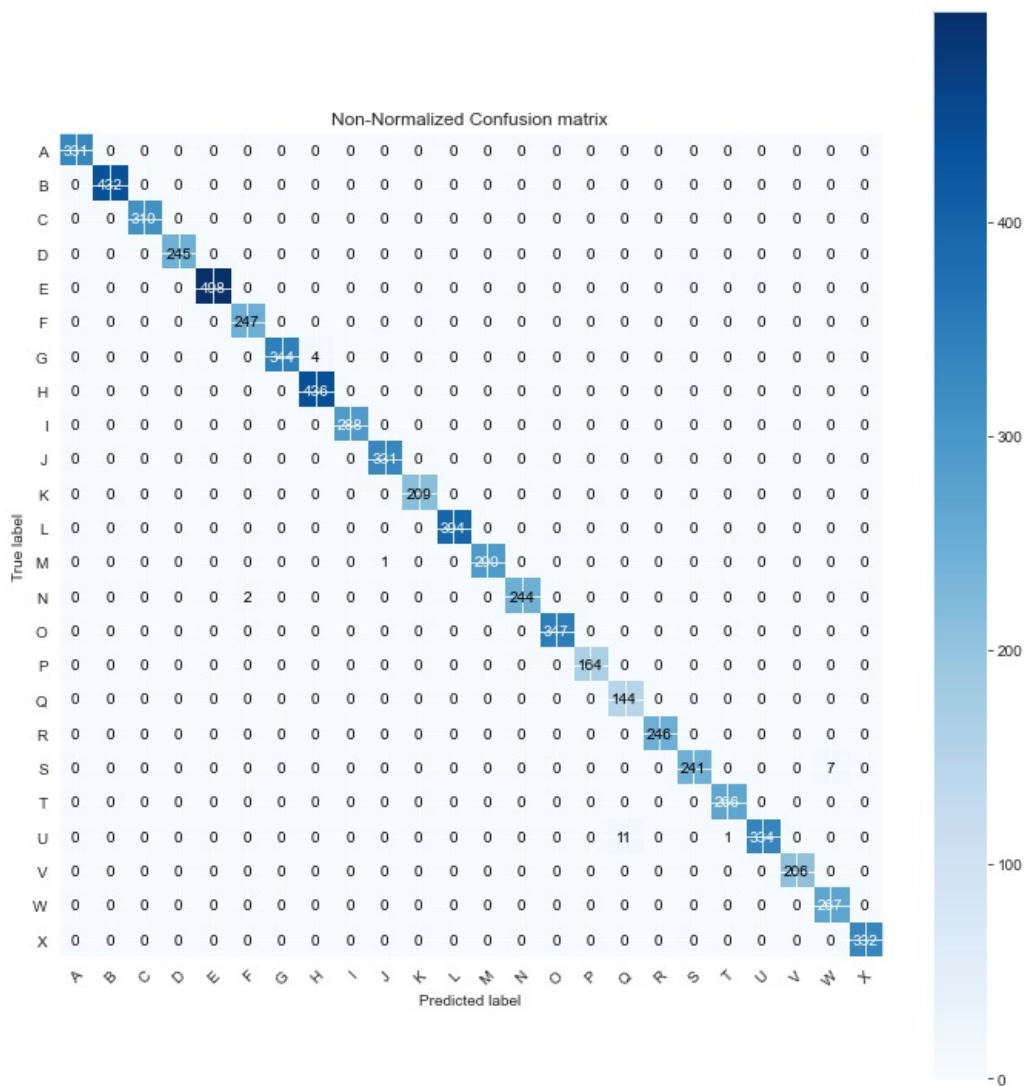
    fig.tight_layout()
    return ax
np.set_printoptions(precision=2)
```

In [26]:

```
#Specifying class labels
class_names = ['A', 'B', 'C', 'D', 'E', 'F', 'G', 'H', 'I', 'J', 'K', 'L', 'M', 'N', 'O', 'P', 'Q', 'R', 'S', 'T',
               'U', 'V', 'W', 'X', 'Y']
```

```
plt.figure(figsize=(20,20))
plot_confusion_matrix(y_test, predicted_classes, classes = class_names, title='Non-Normalized Confusion matrix')
plt.show()
```


<Figure size 1440x1440 with 0 Axes>



```
In [28]:
plt.figure(figsize=(35,35))
plot_confusion_matrix(y_test, predicted_classes, classes = class_names, normalize=True, title='Normalized Confusion matrix')
plt.show()
```

Normalized confusion matrix

<Figure size 2520x2520 with 0 Axes>



In [40]:

```
from sklearn.metrics import confusion_matrix,classification_report,accuracy_score

print(classification_report(y_test, predicted_classes))
print("Accuracy:",accuracy_score(y_test, predicted_classes)*100)
```

	precision	recall	f1-score	support
0.0	1.00	1.00	1.00	331
1.0	1.00	1.00	1.00	432
2.0	1.00	1.00	1.00	310
3.0	1.00	1.00	1.00	245
4.0	1.00	1.00	1.00	498
5.0	0.99	1.00	1.00	247
6.0	1.00	0.99	0.99	348
7.0	0.99	1.00	1.00	436
8.0	1.00	1.00	1.00	288
10.0	1.00	1.00	1.00	331
11.0	1.00	1.00	1.00	209
12.0	1.00	1.00	1.00	394
13.0	1.00	1.00	1.00	291
14.0	1.00	0.99	1.00	246
15.0	1.00	1.00	1.00	347
16.0	1.00	1.00	1.00	164
17.0	0.93	1.00	0.96	144
18.0	1.00	1.00	1.00	246
19.0	1.00	0.97	0.99	248
20.0	1.00	1.00	1.00	266
21.0	1.00	0.97	0.98	346
22.0	1.00	1.00	1.00	206
23.0	0.97	1.00	0.99	267
24.0	1.00	1.00	1.00	332
accuracy			1.00	7172
macro avg	0.99	1.00	1.00	7172
weighted avg	1.00	1.00	1.00	7172

Accuracy: 99.63747908533185

In [30]:

```
predicted_classes.shape
```

Out[30]:

(7172,)

In [31]:

```
y_test.shape
```

Out[31]:

(7172,)

In [42]:

```
#Convert To One Hot encoded Vector
labels = np.zeros((y_test.shape[0],25))

temp = pd.DataFrame(labels)
temp.head()
```

Out[42]:

	0	1	2	3	4	5	6	7	8	9	...	15	16	17	18	19	20	21	22	23	24
0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
1	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
2	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
3	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
4	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0

5 rows × 25 columns

In [43]:

```
for i in range(y_test.shape[0]):
    index = int(y_test[i])
    labels[i][index] = 1
```

In [44]:

```
print(labels.shape)
```

(7172, 25)

In [45]:

```
temp = pd.DataFrame(y_test)
temp.head()
```

Out[45]:

	0
0	6.0
1	5.0
2	10.0
3	0.0
4	3.0

In [46]:

```
temp = pd.DataFrame(labels)
temp.head()
```

Out[46]:

	0	1	2	3	4	5	6	7	8	9	...	15	16	17	18	19	20	21	22	23	24
0	0.0	0.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
1	0.0	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
2	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
3	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
4	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0

5 rows × 25 columns

In [47]:

```
def convert_to_one_hot_encoded(y):
    output_classes = 25
    labels = np.zeros((y.shape[0],output_classes))
    for i in range(y.shape[0]):
        index = int(y[i])
        labels[i][index] = 1
    return labels
```

In [48]:

```
from sklearn.metrics import roc_auc_score

predictions = convert_to_one_hot_encoded(predicted_classes)
y_actual = convert_to_one_hot_encoded(y_test)

scores = []
for i in range(y_actual.shape[0]):
    scores.append(roc_auc_score(y_actual[i],predictions[i]))

scores = np.array(scores)
print(np.mean(scores)*100)
```

99.81118702361034

In []: